

EXPLAINABLE AI HIRING OS • WHITE, NOT BLACK

WHITEBOX AI

Not a faster black box — **a white one you can see into.**

Every score explained. Every decision human. Every action auditable.



RANKED DRIVERS

EVIDENCE CHIPS

REASONING TRACE

HUMAN DECISION

AUDIT LOG

Hiring AI became a black box nobody can question.

Companies drowning in applications let AI/ATS decide who ever reaches a human. The candidate gets silence, the recruiter gets a number, and hidden bias gets **replicated instead of corrected**.

Human presence alone doesn't fix it — research shows people follow biased recommendations. Visibility and control do.

1000s

of applications per role —
AI screens them before any
human

0

reasons shown to a
rejected candidate — or to
the recruiter

≠

different LLMs rank
identical résumés
differently

**HIGH-
RISK**

how the EU AI Act classifies
hiring AI. NYC already
requires bias audits

Three curves crossed in the last 24 months.

01 • ADOPTION

AI screening became default

After LLMs, resume ranking and AI interviews went from novelty to standard pipeline — faster than anyone audited them.

02 • REGULATION

Law caught up

EU AI Act puts hiring AI in the high-risk class. NYC Local Law 144 already mandates bias audits. Explainability is becoming a legal requirement.

03 • TRUST

Both sides stopped believing

Candidates mass-apply with AI; employers auto-reject with AI. Nobody can explain an outcome — and everybody needs to.

The reason-code playbook — legally proven in credit scoring for 40 years — has never been applied to hiring. Until now.

Every score arrives fully explained.

WhiteBox AI is a two-sided screening platform. The AI **structures information**; the human **decides**. That contract is enforced by the product, not by policy.



01 Ranked drivers

Top factors that pushed the score up or down — the FICO reason-code model, the only legally-tested form of explainable adverse action.

02 Confidence as a phrase

"High confidence — consistent evidence across three answers", never a bare 0.87.

03 Evidence chips

Every claim links to the actual interview fragment it came from.


04 Reasoning trace

The full collapsible path from answer to score — inspectable, contestable.


05 Mandatory human exit + audit log

Accept / change / recompute / reject-with-reason — written to an append-only, hash-chained audit log.

A dense instrument for people who hire at volume.




CANDIDATE RANKING — TIERS, DRIVERS, BULK ACTIONS




SCORECARD — RANKED DRIVERS • CONFIDENCE • EVIDENCE • TRACE
• HUMAN EXIT

Calm, transparent, dignified.




STRUCTURED AI INTERVIEW — ADAPTIVE FOLLOW-UPS, GENTLE PROGRESS



STATUS — CLEAR STAGE, HONEST COPY, THE WHITE-BOX TRUST PANEL

Six steps. One auditable pipeline.



Bias monitoring (Fairlearn/Aequitas) and GDPR-grade data lifecycle run underneath the whole pipeline.

Compliance turns a nice-to-have into a procurement requirement.

\$3B+

AI recruitment software market, growing with every enterprise AI-screening rollout

ICP

High-volume enterprise hirers

Banks, telecom, retail — hundreds to thousands of applications per role, existing ATS, real regulatory exposure. Buyer: HR director / head of TA. Payer: the company. Users: recruiters — and candidates, free.

ENTRY PATH

KZ / Central Asia → EU / US

Design partners at home first; then the compliance-driven mid-market as EU AI Act enforcement matures. Global SaaS from day one.

Everyone scores. Nobody explains.

CAPABILITY	AI INTERVIEWERS HIREVUE • MERCOR • BRIGHTHIRE	ATS SCORECARDS ASHBY • GREENHOUSE	WHITEBOX AI
Per-competency scoring, structured interviews	✓	✓	✓
Ranked score drivers (reason codes)	—	—	✓
Evidence linked to every score	—	—	✓
Human final decision enforced in-product	—	partial	✓
Append-only, hash-chained audit log	—	—	✓

Explainability is our architecture, not a feature toggle — black-box competitors can't retrofit it without rebuilding.

Land on one vacancy. Expand to the org.

MODEL

B2B SaaS

Pilot: per-vacancy pricing on one live role — prove time-to-shortlist and auditability.

Expand: seat-based annual contracts for the TA team. The company pays; candidates never do.

GO-TO-MARKET

Founder-led, compliance-driven

- 1 — Design partners via warm intros, KZ / Central Asia.
- 2 — Direct outbound to HR directors of high-volume employers.
- 3 — Content on explainable & compliant AI hiring; ATS integrations (Greenhouse, Ashby) as distribution.

MOAT

Tech + data

Architecture built around contestable decisions — hard to retrofit.

Growing dataset of **explained, human-validated** hiring decisions per competency; compliance depth becomes a regulatory moat as enforcement matures.

SALES CYCLE: 3-6 MONTHS • CHANNELS: DIRECT SALES • CONTENT • PARTNERSHIPS • INTEGRATIONS/API

The MVP already exists.

You just saw it.

- ✓
P0

Foundation — design tokens, app shell, 3K palette, white-box primitive
- ✓
P1

Core differentiator — ranking, explainable scorecard, HITL + hash-chained audit
- P2

Configuration — competency builder, onboarding, full interview flow
- P3

Real AI pipeline — transcription, scoring, bias monitoring (FastAPI · LangGraph · Temporal)
- P4

Polish — 3D cube hero, motion system, hardening

GITHUB.COM/PIP00SYA/WHITEBOX-AI • NEXT.JS 15 • TYPESCRIPT •
TOKEN-DRIVEN DESIGN SYSTEM



LIVE PROTOTYPE — HR DASHBOARD

Two founders.

Full stack of the problem.

FOUNDER • CEO

Nurkhan Aimukatov

Product, design system and engineering.
Authored the full product specification — vision, HR and candidate specs, build sequence — and built the working MVP end-to-end (Next.js, TypeScript) with an AI-assisted development workflow.

CO-FOUNDER

Business development

Partnerships and go-to-market: design-partner pipeline in Kazakhstan and Central Asia, first pilot deployments with high-volume employers.

HIRING NEXT: ML/AI ENGINEER (MULTIMODAL ANALYSIS) • B2B SALES LEAD • COMPLIANCE / ORG-PSYCHOLOGY ADVISOR

Raising **\$100K** pre-seed for a 12-month runway.

USE OF FUNDS

Product & AI engineering — vacancy wizard, competency builder, full multimodal interview pipeline

LLM inference & infrastructure

Compliance & security — audit, bias testing, data protection

First pilots — design-partner onboarding

3 MONTHS

Build sequence complete (P2–P4)

Closed beta with **2–3 design partners**

6 MONTHS

5–10 pilot vacancies screened end-to-end

Measured time-to-shortlist & override rates

First **paid pilots**